

Homework 2

Aaron Politsky

October 9, 2015

5.1

Load Libraries and Source Dependencies

```
library(data.table)
library(ggplot2)
library(kknn)
source("docv.R")
```

Import the Used Car data

```
# download data and read data into data.table format
if(!exists("used_cars")) {
  used_cars <- fread(
    'https://raw.githubusercontent.com/ChicagoBoothML/DATA__UsedCars/master/UsedCars_small.csv')
}

# sort data set by increasing mileage
setkey(used_cars, mileage)
```

Cross-validate using 5-fold, for k from 2 to 100

```
n.folds <- 5

set.seed(99) # to be consistent with hw1
kv <- 2:100
cv <- docvknn(x=matrix(used_cars$mileage),
              y=used_cars$price,
              k=kv,
              nfold=n.folds)
```

```
## in docv: nset,n,nfold: 99 1000 5
## on fold: 1 , range: 1 : 200
## on fold: 2 , range: 201 : 400
## on fold: 3 , range: 401 : 600
## on fold: 4 , range: 601 : 800
## on fold: 5 , range: 801 : 1000
```

```

# convert to RMSE
cv <- sqrt(cv/length(used_cars$price))

# Choose the k with the minimum Cross-validation RMSE
k.cv <- which.min(cv)

```

We will use a k value of $k = 19$. Let's fit using this value and the eyeball value of 30:

```

# Produce a model with this k
cv.model <- kknnp(price ~ mileage,
                  train=used_cars, test=used_cars[, .(mileage)],
                  k=k.cv, kernel='rectangular')


# add the predicted values to our data.table
used_cars$cv.predicted <- cv.model$fitted.values

# Produce the eyeball fit from hw1
k <- 30
eyeball <- kknnp(price ~ mileage,
                  train=used_cars, test=used_cars[, .(mileage)],
                  k=k, kernel='rectangular')

# add this prediction to used_cars data.table
used_cars$eyeball.predicted <- eyeball$fitted.values

g <- ggplot(used_cars) +
  geom_point(aes(x=mileage, y=price, color='actual'), size=1) +
  ggtitle('Used Cars: price vs. mileage') +
  xlab('mileage') + ylab('price')
g <- g +
  geom_line(aes(x=mileage, y=cv.predicted, color='cv predicted'), size=0.6) +
  geom_line(aes(x=mileage, y=eyeball.predicted, color='eyeball predicted'), size=0.6) +
  scale_colour_manual(name='price',
                      values=c(actual='blue',
                                "cv predicted"='red',
                                "eyeball predicted"='green'))
plot(g)

```



In this case, our eyes have not deceived us. The fits are very close. This makes sense, because $cv[19] = 9248.2401549$ and $cv[30] = 9284.1316686$. The fit having the larger k value seems to have slightly less variance (the green curve seems to be “within” the red curve) than the fit with the smaller k value, which is consistent with the theory of knn.

Predicting the price of a car with 100k miles

Now, let’s predict the price of a car having 100k miles using our cross-validation k to train a model using the entire used car dataset.

```
# simulate a car with 100k miles
car.w.100k.miles <- data.frame(mileage=100000)
predicted.price.car.w.100k.miles <-
  knn(price ~ mileage,
      train=used_cars, test=car.w.100k.miles,
      k=k.cv, kernel='rectangular')$fitted.values
predicted.price.car.w.100k.miles
```

```
## [1] 18662.63
```